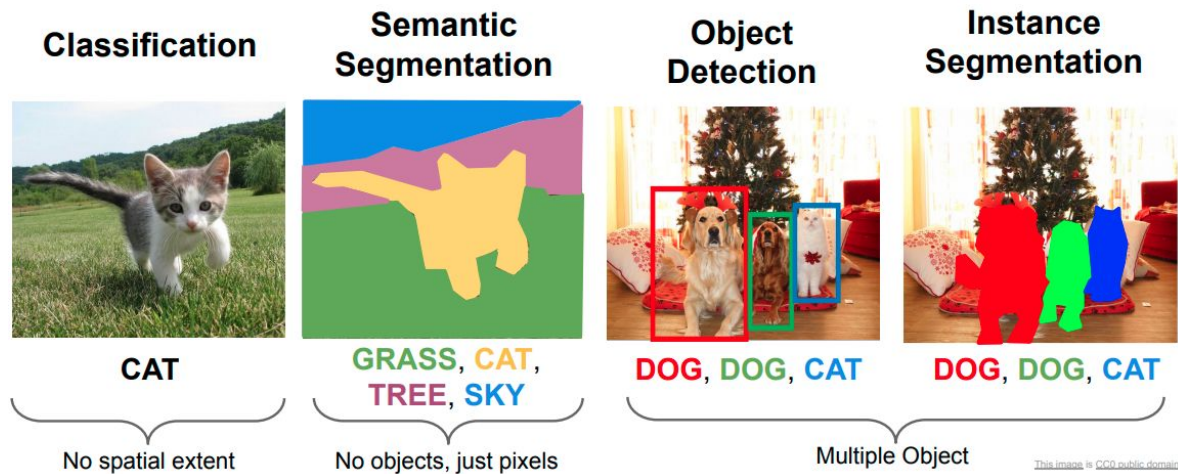


Problem Statement




This image is CC0 public domain

Image Source: <http://cs231n.stanford.edu/>

Object Localization



Object Classification is the task of identifying that picture is a dog



Object Localization involves the class label as well as a bounding box to show where the object is located.

Image Source:

<https://www.commonlounge.com/discussion/c9975025c9ff473c8f9ed2c4b1c3ea6a#image-localization-detection-segmentation>




Image Source: [MNIST Handwritten Digits Classification using a Convolutional Neural Network \(CNN\)](#)




Image Source: deeplearning.ai

Idea: Sliding Window




Image Source: <https://datahacker.rs/deep-learning-object-detection/>

Idea: Region Proposal

Felzenszwalb's Algorithm [2004] :
Efficient Graph-Based Image Segmentation

- Turn the image into an undirected graph $G = (V, E)$
 - Vertex $v_i \in V$ is a pixel
 - Edge $e = (v_i, v_j) \in E$ connects two vertices
 - Weight $w(v_i, v_j)$ measure dissimilarity between v_i and v_j
- A segmentation solution S is a partition of V into connected components $\{C\}$




Image Source: [Felzenszwalb, P.F., Huttenlocher, D.P. Efficient Graph-Based Image Segmentation. International Journal of Computer Vision 59, 167–181 \(2004\).](#)

Idea: Region Proposal

Selective Search :

“Class Agnostic Object Detector”

- Start from Felzenszwalb’s Segmentations
- Recursively combine similar regions into larger ones
- Convert regions to boxes



Image Source: http://vision.stanford.edu/teaching/cs231b_spring1415

Overview: R-CNNs

R-CNN

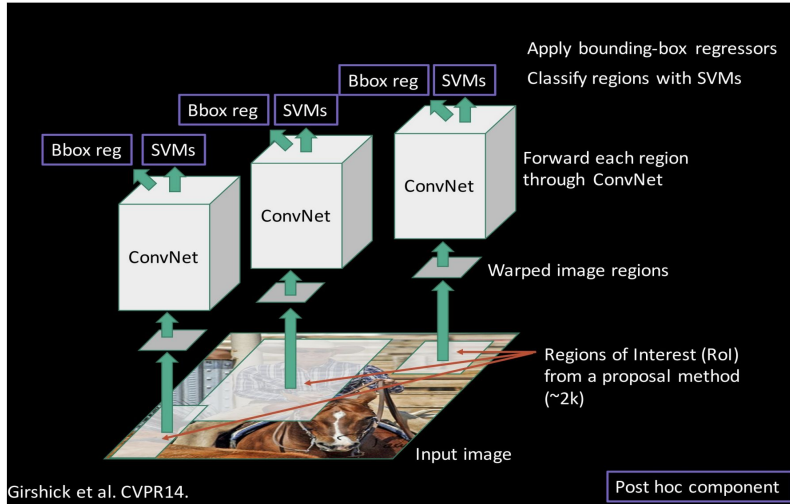


Image: Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014

Fast R-CNN

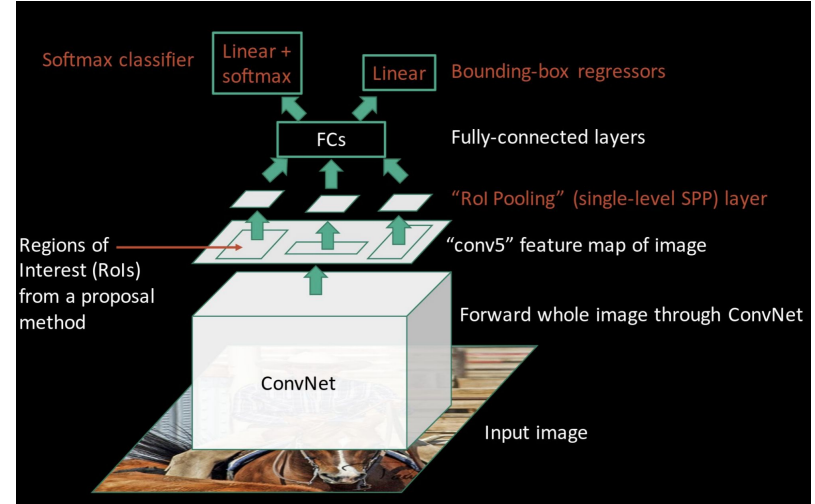



Image: Girshick, "Fast R-CNN", ICCV 2015

Overview: R-CNNs

Faster R-CNN



Mask R-CNN




Image: Ren et al, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks", NIPS 2015

Image Source: [Object Detection for Dummies Part 3: R-CNN Family](#)

Model: Mask R-CNN

Transfer Learning:

- Start with a pretrained model
- Replace classifier and mask predictor
- Fine tune

Loss Function: $\mathcal{L} = \mathcal{L}_{\text{cls}} + \mathcal{L}_{\text{box}} + \mathcal{L}_{\text{mask}}$

where;

$$\mathcal{L}(\{p_i\}, \{t_i\}) = \frac{1}{N_{\text{cls}}} \sum_i \mathcal{L}_{\text{cls}}(p_i, p_i^*) + \frac{\lambda}{N_{\text{box}}} \sum_i p_i^* \cdot L_1^{\text{smooth}}(t_i - t_i^*)$$

$$\mathcal{L}_{\text{cls}}(p_i, p_i^*) = -p_i^* \log p_i - (1 - p_i^*) \log(1 - p_i)$$

$$\mathcal{L}_{\text{mask}} = -\frac{1}{m^2} \sum_{1 \leq i, j \leq m} [y_{ij} \log \hat{y}_{ij}^k + (1 - y_{ij}) \log(1 - \hat{y}_{ij}^k)]$$





Image Source: [Object Detection for Dummies Part 3: R-CNN Family](#)

Model: Mask R-CNN

Transfer Learning:

- Start with a pretrained model
- Replace classifier and mask predictor
- Fine tune



Loss Function: $\mathcal{L} = \mathcal{L}_{\text{cls}} + \mathcal{L}_{\text{box}} + \mathcal{L}_{\text{mask}}$

Image Source: [Object Detection for Dummies Part 3: R-CNN Family](#)

where;

$$\mathcal{L}(\{p_i\}, \{t_i\}) = \frac{1}{N_{\text{cls}}} \sum_i \mathcal{L}_{\text{cls}}(p_i, p_i^*) + \frac{\lambda}{N_{\text{box}}} \sum_i p_i^* \cdot L_1^{\text{smooth}}(t_i - t_i^*)$$

$$\mathcal{L}_{\text{cls}}(p_i, p_i^*) = -p_i^* \log p_i - (1 - p_i^*) \log(1 - p_i)$$

$$\mathcal{L}_{\text{mask}} = -\frac{1}{m^2} \sum_{1 \leq i, j \leq m} [y_{ij} \log \hat{y}_{ij}^k + (1 - y_{ij}) \log(1 - \hat{y}_{ij}^k)]$$

Tutorials:

PyTorch:

- TORCHVISION OBJECT DETECTION FINETUNING TUTORIAL
https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html#torchvision-object-detection-finetuning-tutorial

Tensorflow:

- Object Detection API with Tensorflow 2
https://github.com/tensorflow/models/blob/master/research/object_detection/g3doc/tf2.md

Keras:

- Object Detection with RetinaNet
<https://keras.io/examples/vision/retinanet/>

Results: Metrics

- **Precision** = $TP / (TP + FP)$
- **Recall** = $TP / (TP + FN)$
- **Average Precision (AP):**
Area under Precision/Recall curve
- **Mean AP (mAP):**
Average AP over all classes

mAP@0.5 → TP if IoU > 0.5

